


Finetuning LLM(Mistral-7B) on Colab

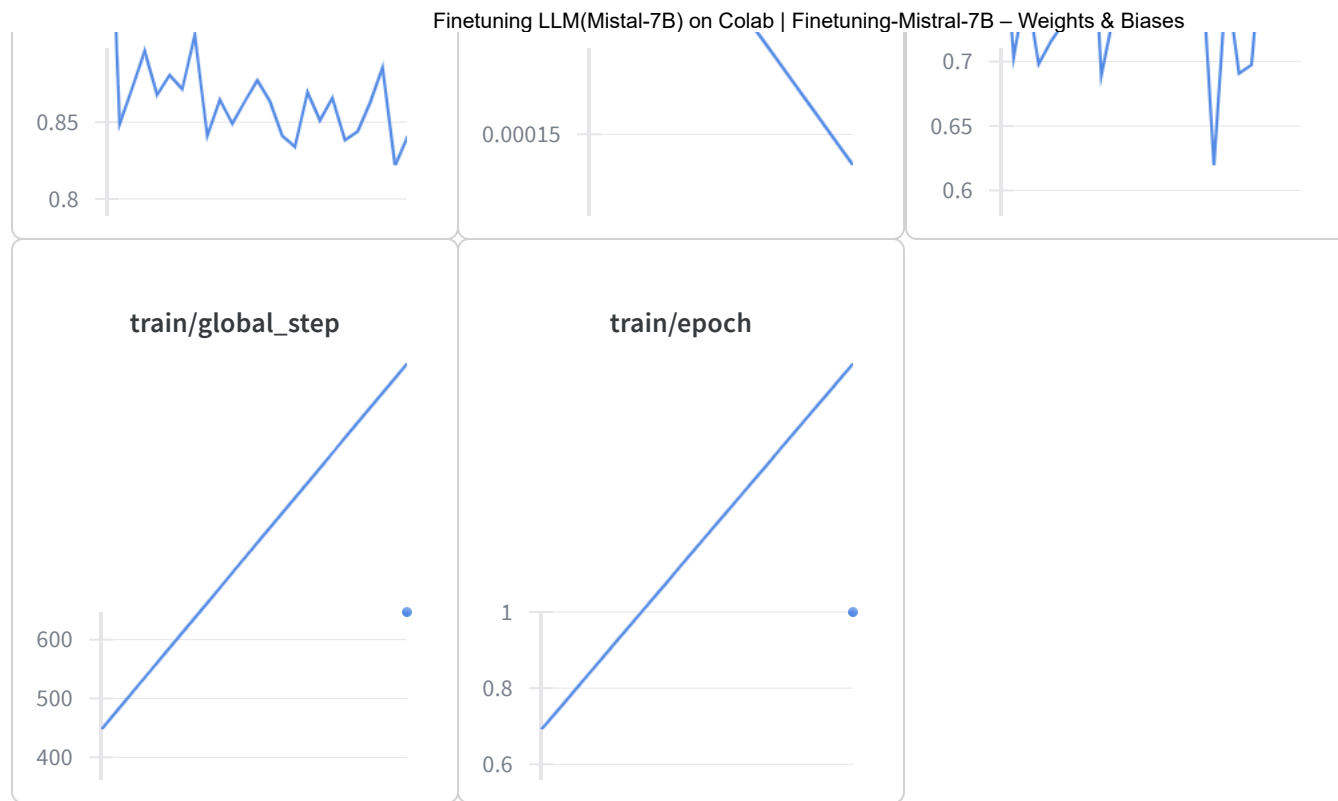
This project fine-tunes the Mistral-7B-Instruct model on Colab free GPUs using QLoRA (4-bit quantization) and LoRA adapters for memory-efficient instruction tuning. It leverages the Unsloth library to load the quantized model, apply LoRA to attention and feed-forward layers, and train with the Alpaca-cleaned dataset (formatted in an instruction–input–response template). Training is performed with mixed precision, an 8-bit AdamW optimizer, gradient accumulation, and gradient checkpointing to fit

Abdul Basit

Created on September 28 | Last edited on September 28

✓ Section 1:





Created with ❤️ on Weights & Biases.

<https://wandb.ai/basitmal36-university-of-paris-saclay/Finetuning-Mistral-7B/reports/Finetuning-LLM-Mistal-7B-on-Colab--VmIldzoxNDU1ODc2Nw>